



# Ceph RGW Cache Prefetching for Batch Jobs

Xun Lin  
Yang Qiao  
Tianyi Tang  
Gang Wei  
Zhangyu Wan



# Progress

- Problems in the last Demo:
  - DAG is only extracted at application level
  - Timing information is not clear for running Dijkstra



# Progress

DAG:

- Change the Spark source code, compile it, extract DAG before actually running the jobs (DAG is extracted at platform level)
- Send DAG information to Kariz by HTTP request

Kariz:

- Correct DAG string convert function
- Run a test with Kariz server, cache simulator and DAG simulator

# Taiga

DEMO 4 BU-CEPH RGW CACHE PREFETCHIN... 25 OCT 2019-31 OCT 2019



100%

5 total points


5 completed points

0 open tasks

0 closed tasks



0 iocaine doses



# Taiga

DEMO 4 - BU-CEPH RGW CACHE PREFETCHIN... 01 NOV 2019-07 NOV 2019



60%

25 total points


15 completed points

1 open tasks

4 closed tasks



0 iocaine doses





## Spark Integration with Kariz

- Why?

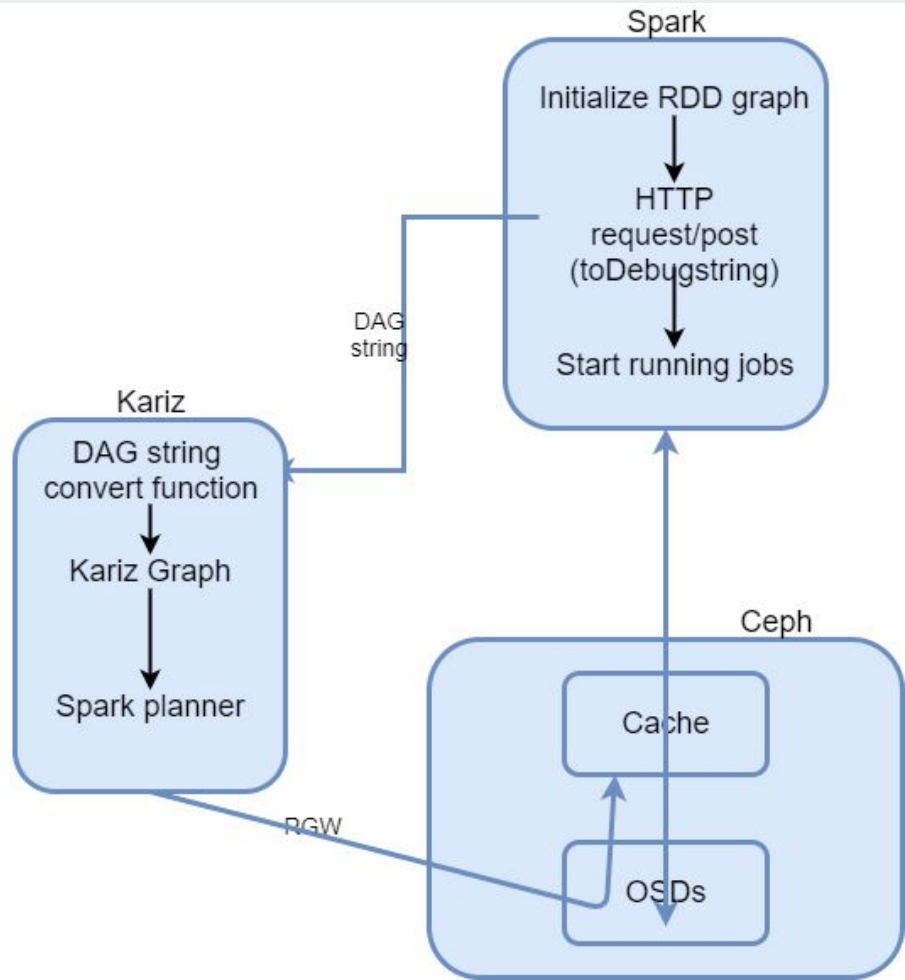
Kariz need DAG information to determine what and when to prefetch.

- How?

Send HTTP **POST** request with DAG string.

# How system works

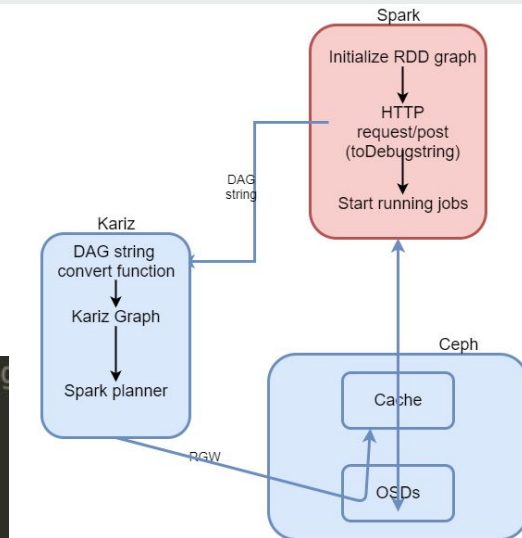
1. Spark generates the DAG before running jobs.
2. Spark sends DAG string to Kariz.
3. Kariz translates the string into its Graph class and submit the json file to its planner.
4. The planner prefetch data from Ceph into D3N cache
5. Spark runs the jobs ( files are already in cache)



# RDD.scala

(spark/core/src/main/scala/org/apache/spark/rdd/RDD.scala)

```
1829  /** A description of this RDD and its recursive dependencies for debugg
1830  def toDebugString: String = {
1831    // Get a debug description of an rdd without its children
1832    def debugSelf(rdd: RDD[_]): Seq[String] = {
1833      import Utils.bytesToString
1834
1835      val persistence = if (storageLevel != StorageLevel.NONE) storageLevel.description else ""
1836      val storageInfo = rdd.context.getRDDStorageInfo(_.id == rdd.id).map(info =>
1837        "    CachedPartitions: %d; MemorySize: %s; ExternalBlockStoreSize: %s; DiskSize: %s".format(
1838          info.numCachedPartitions, bytesToString(info.memSize),
1839          bytesToString(info.externalBlockStoreSize), bytesToString(info.diskSize))
1840
1841      s"$rdd [$persistence]" +: storageInfo
1842    }
1843
```

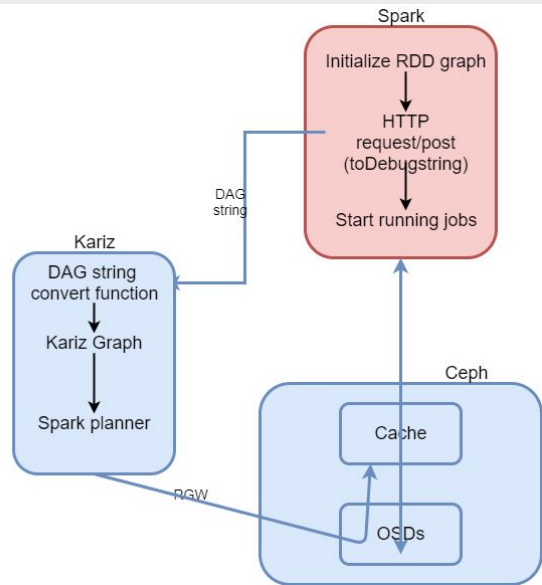




# SparkContext.scala


(spark/core/src/main/scala/org/apache/spark/SparkContext.scala)

```
2091 def runJob[T, U: ClassTag](
2092     rdd: RDD[T],
2093     func: (TaskContext, Iterator[T]) => U,
2094     partitions: Seq[Int],
2095     resultHandler: (Int, U) => Unit): Unit = {
2096     if (stopped.get()) {
2097         throw new IllegalStateException("SparkContext has been shutdown")
2098     }
2099     val callSite = getCallSite
2100     val cleanedFunc = clean(func)
2101     logInfo("Starting job: " + callSite.shortForm)
2102     logInfo("RDD's recursive dependencies:\n" + rdd.toDebugString)
2103     logInfo("Sending HTTP Request to Kariz")
2104     val url = "http://kariz-1:3188/api/newstage"
2105     val post = new HttpPost(url)
2106     post.addHeader("Content-Type", "text/plain")
2107     post.addHeader("Accept", "text/plain")
2108     val client = new DefaultHttpClient
2109     val nameValuePairs = new ArrayList<NameValuePair>(1)
2110     nameValuePairs.add(new BasicNameValuePair("dag", rdd.toDebugString))
2111     post.setEntity(new UrlEncodedFormEntity(nameValuePairs))
2112     // send the post request
2113     val response = client.execute(post)
2114     dagScheduler.runJob(rdd, cleanedFunc, partitions, callSite, resultHandler, localProperties.get)
2115     progressBar.foreach(_.finishAll())
2116     rdd.doCheckpoint()
2117 }
```



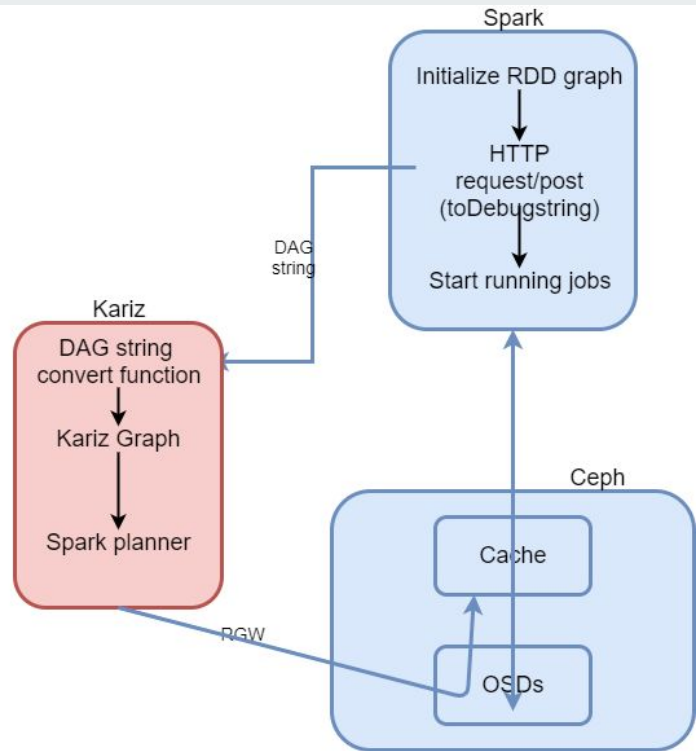


# Demo



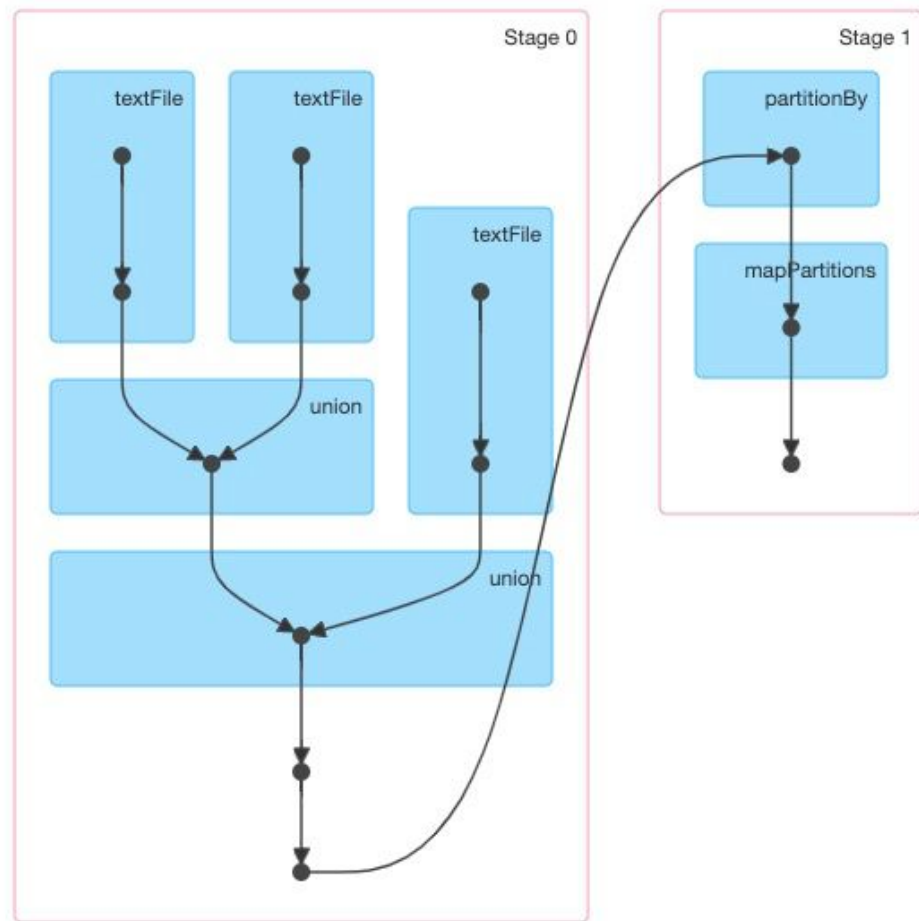
## DAG string convert function

- Previously we thought the input of one node is the output rdd of the following node.
- From the Spark UI: the input of one node is all the output rdds below the node whose rdd number is smaller than the node output rdd number.

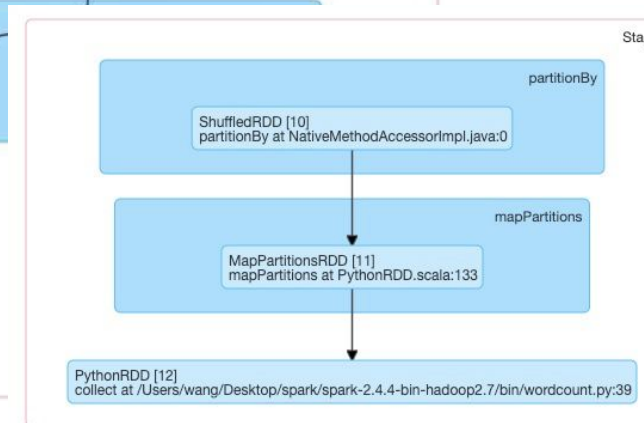
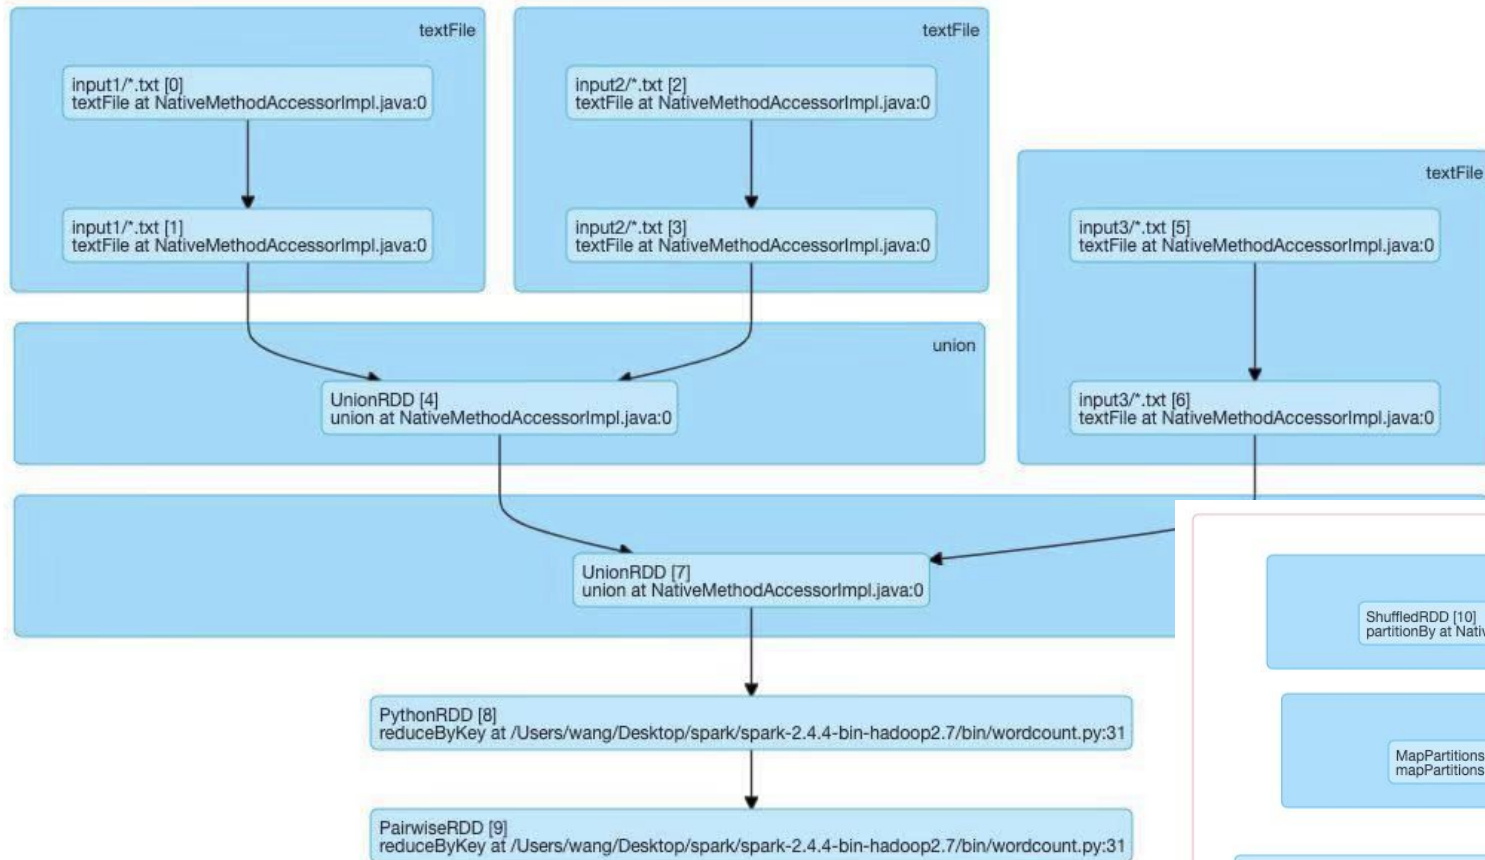


# DAG visualization (multiple path)

```
(8) PythonRDD[12] at collect at
/Users/joe/Desktop/Spark_Source/spark/bin/wordcount.py:39 []
| MapPartitionsRDD[11] at mapPartitions at PythonRDD.scala:133 []
| ShuffledRDD[10] at partitionBy at NativeMethodAccessorImpl.java:0 []
+-(8) PairwiseRDD[9] at reduceByKey at
/Users/joe/Desktop/Spark_Source/spark/bin/wordcount.py:31 []
| PythonRDD[8] at reduceByKey at
/Users/joe/Desktop/Spark_Source/spark/bin/wordcount.py:31 []
| UnionRDD[7] at union at NativeMethodAccessorImpl.java:0 []
| UnionRDD[4] at union at NativeMethodAccessorImpl.java:0 []
| input1/*.txt MapPartitionsRDD[1] at textFile at NativeMethodAccessorImpl.java:0 []
| input1/*.txt HadoopRDD[0] at textFile at NativeMethodAccessorImpl.java:0 []
| input2/*.txt MapPartitionsRDD[3] at textFile at NativeMethodAccessorImpl.java:0 []
| input2/*.txt HadoopRDD[2] at textFile at NativeMethodAccessorImpl.java:0 []
| input3/*.txt MapPartitionsRDD[6] at textFile at NativeMethodAccessorImpl.java:0 []
| input3/*.txt HadoopRDD[5] at textFile at NativeMethodAccessorImpl.java:0 []
```



Stage 0



# Longest Path

- Why need longest path ?

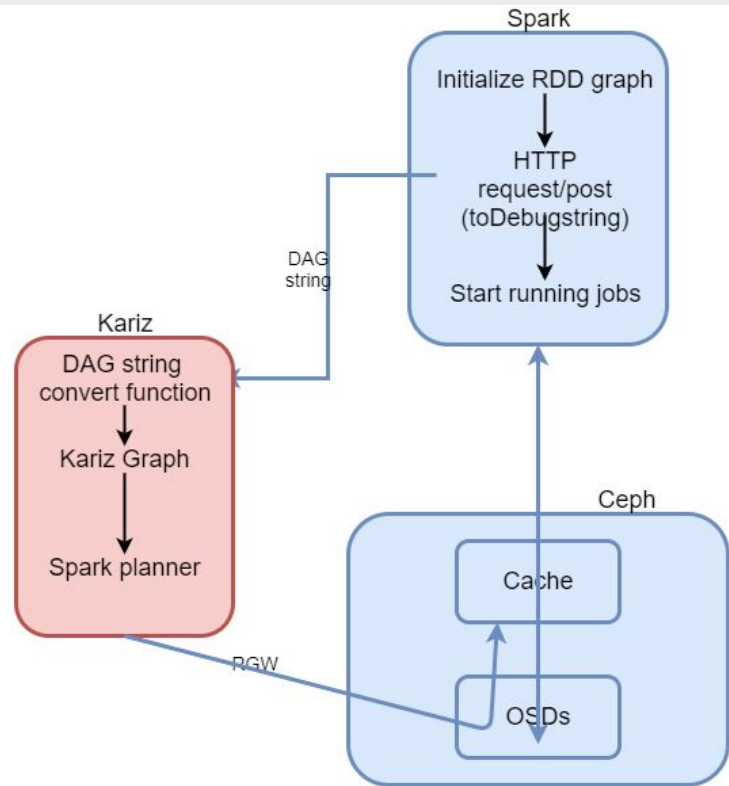
Kariz is designed to prefetch the data that **reduce runtime the most**

How many cache we have

Statistics ,( job types) from spark → prediction of runtime

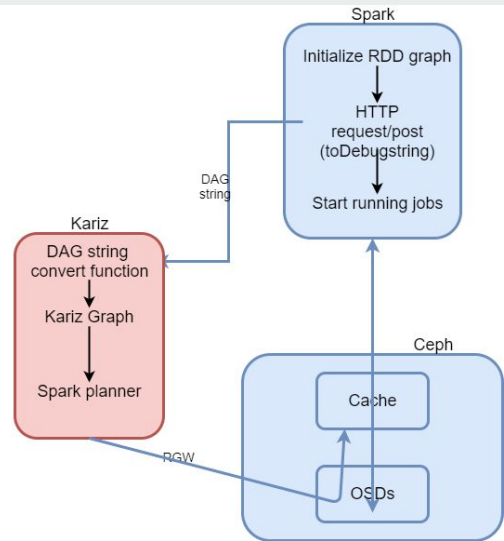
Use the longest path → cache planner for spark

Analysis: Worst case/ Best case/ Average case



# Longest Path

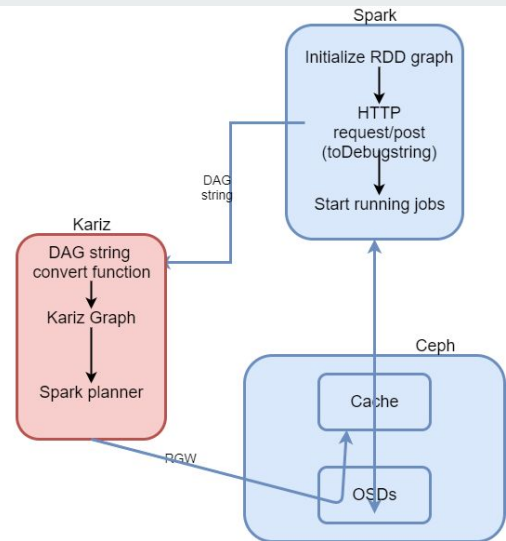
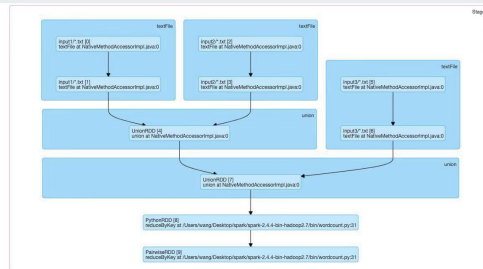
- Timing information of all nodes is currently randomly set.
- Individual job execution times can be predicted with data from past executions of the same job type and network and storage system bandwidth, which is measured by Kariz.
- We will make Spark send job type to Kariz and integrate the longest path into Spark planner in the next sprint.



# Longest Path

{0: {'output': 'HadoopRDD[5]', 'inputs': ['input3/\*.txt']}, 1: {'output': 'MapPartitionsRDD[6]', 'inputs': ['HadoopRDD[5]', 'input3/\*.txt']}, 2: {'output': 'HadoopRDD[2]', 'inputs': ['input2/\*.txt']}, 3: {'output': 'MapPartitionsRDD[3]', 'inputs': ['HadoopRDD[2]', 'input2/\*.txt']}, 4: {'output': 'HadoopRDD[0]', 'inputs': ['input1/\*.txt']}, 5: {'output': 'MapPartitionsRDD[1]', 'inputs': ['HadoopRDD[0]', 'input1/\*.txt']}, 6: {'output': 'UnionRDD[4]', 'inputs': ['MapPartitionsRDD[3]', 'MapPartitionsRDD[1]']}, 7: {'output': 'UnionRDD[7]', 'inputs': ['MapPartitionsRDD[6]', 'UnionRDD[4]']}, 8: {'output': 'PythonRDD[8]', 'inputs': ['UnionRDD[7]']}, 9: {'output': 'PairwiseRDD[9]', 'inputs': ['PythonRDD[8]']}, 10: {'output': 'ShuffledRDD[10]', 'inputs': ['PairwiseRDD[9]']}, 11: {'output': 'MapPartitionsRDD[11]', 'inputs': ['ShuffledRDD[10]']}, 12: {'output': 'PythonRDD[12]', 'inputs': ['MapPartitionsRDD[11]']}

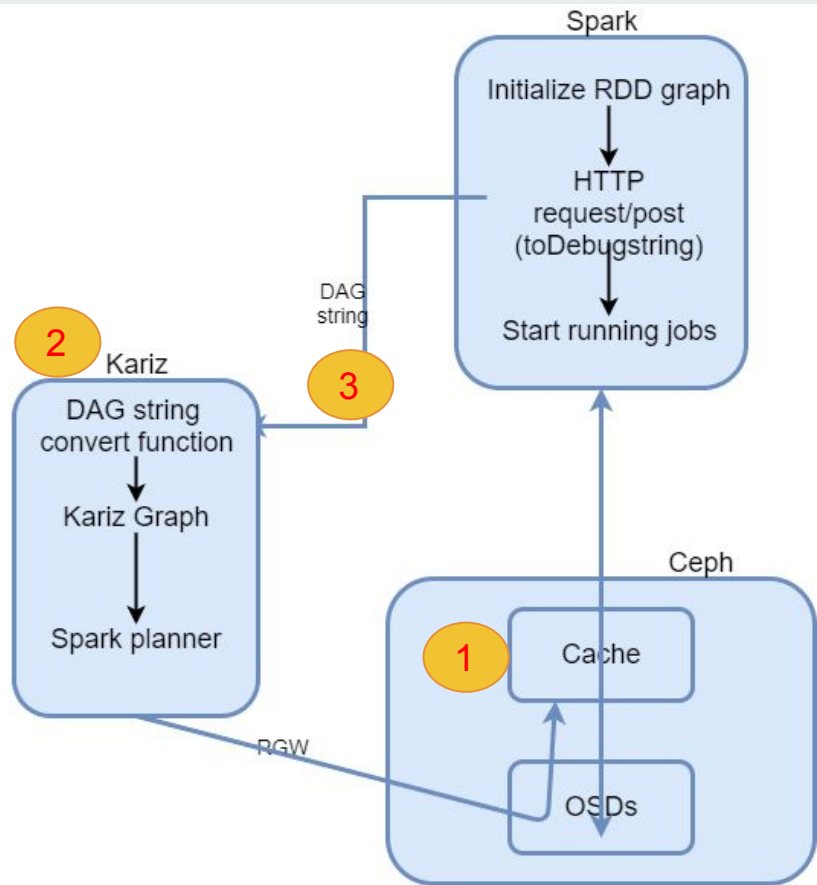
Longest path starting at node 0 ends at node 12 with length -7  
 The longest route is [0, 1, 7, 8, 9, 10, 11, 12]  
 Longest path starting at node 2 ends at node 12 with length -8  
 The longest route is [2, 3, 6, 7, 8, 9, 10, 11, 12]  
 Longest path starting at node 4 ends at node 12 with length -8  
 The longest route is [4, 5, 6, 7, 8, 9, 10, 11, 12]





# Kariz simulation

1. Cache simulator
  - Simulates as D3N cache
2. Server simulator with Kariz running on it
3. DAG simulator
  - Submits DAG to Kariz



# Kariz simulation

```
dhcp-acadmin-204-8-153-27:kariz gangwei$ python3 ./api/server.py
* Serving Flask app "server" (lazy loading)
* Environment: production
WARNING: This is a development server. Do not use it in a production deployment.
Use a production WSGI server instead.
* Debug mode: on
* Running on http://0.0.0.0:3188/ (Press CTRL+C to quit)
* Restarting with stat
* Debugger is active!
* Debugger PIN: 141-252-419
127.0.0.1 - - [06/Nov/2019 14:17:13] "POST /api/newdag HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:17:13] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:17:18] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:17:38] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:17:38] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:17:38] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:17:58] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:17:58] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:18:18] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:18:18] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:18:38] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:18:38] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:18:38] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:18:58] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:18:58] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:19:18] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:19:18] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:19:38] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:19:38] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:19:58] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:19:58] "POST /api/newstage HTTP/1.1" 200 -
127.0.0.1 - - [06/Nov/2019 14:20:18] "POST /api/completed HTTP/1.1" 200 -
```

```
cache — Python • Python server.py — 80x16
GangdeMacBook-Pro:cache gangwei$ python3 server.py
* Serving Flask app "server" (lazy loading)
* Environment: production
WARNING: This is a development server. Do not use it in a production deployment.
Use a production WSGI server instead.
* Debug mode: on
* Running on http://0.0.0.0:3187/ (Press CTRL+C to quit)
* Restarting with stat
* Debugger is active!
* Debugger PIN: 141-252-419

clear_cache done: status: global status: {}
Worker: free space: 550, unpinned space: 550, {}
127.0.0.1 - - [06/Nov/2019 14:20:18] "POST /cache/clearcache HTTP/1.1" 200 -

framework_simulator — -bash — 96x22
GangdeMacBook-Pro:framework_simulator gangwei$ python3 main.py
Schedule stage 0 for execution. Estimated runtime: 0.7054561891443976 , elapsed time: 0.7054561891443976
Schedule stage 1 for execution. Estimated runtime: 0.8280478537564019 , elapsed time: 0.7054561891443976
```



## Next sprint

- Kariz GET request
- Spark send job type to Kariz
- Apply the longest path to Spark planner
- Replace Cache simulator with real Ceph Cache