



Ceph RGW Cache Prefetching for Batch Jobs

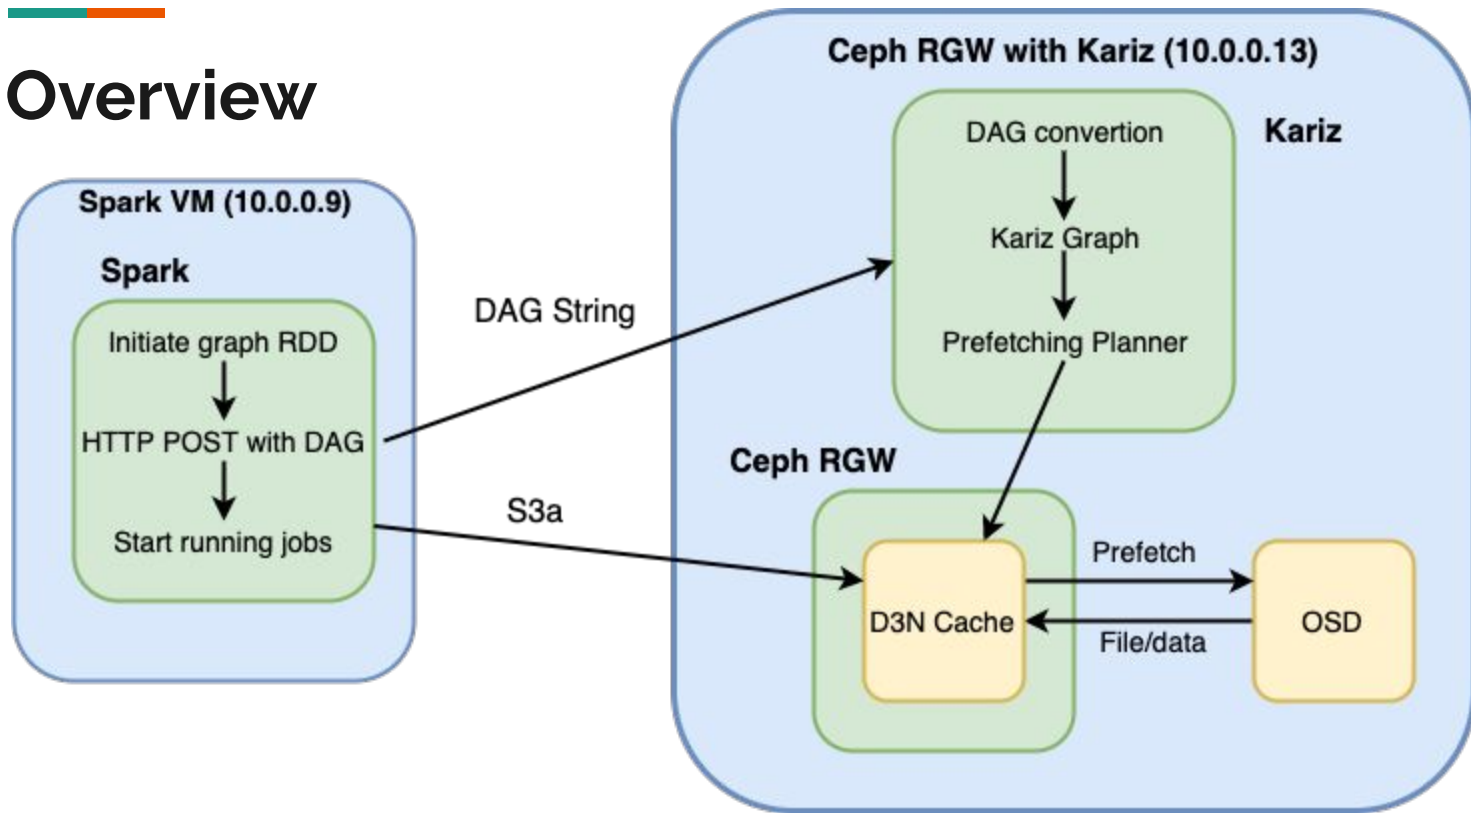
Xun Lin
Yang Qiao
Tianyi Tang
Gang Wei
Zhangyu Wan



Overview

- MVP:
 - Extract DAG out of Spark Applications
 - Find the job dependency path, generate cache planner
 - Prefetch Files/data while running batch jobs
 - Performance evaluation (with & without prefetching)

Overview





Progress

- **Finish End to End implementation**
 - Spark
 - Integration of Spark, Ceph and Kariz
 - Kariz cache planner
 - Parse requests to graph and planner
 - Prefetching
 - Prefetch files/data based on planner

Taiga

DEMO 5 BU-CEPH RGW CACHE PREFET... 08 NOV 2019-21 NOV 2019



78%

45 total points


35 completed points

1 open tasks

5 closed tasks



0 iocaine doses





System Integration

- Environment
 - Spark/Hive in 10.0.0.9
 - Kariz/Ceph RGW in 10.0.0.13
- Communication
 - Spark - Kariz
 - Spark - Ceph RGW
 - Kariz - Ceph RGW



Hive with Spark

1. Start Spark in standalone mode and set environment variable
2. Link Scala and Spark jars in Hive lib folder

```
cd $HIVE_HOME/lib  
ln -s $SPARK_HOME/jars/scala-library*.jar  
ln -s $SPARK_HOME/jars/spark-core*.jar  
ln -s $SPARK_HOME/jars/spark-network-common*.jar
```



Hive with Spark

3. Set the hive.execution.engine to Spark in hive-site.xml (configuration file)

4. Set Spark parameters in hive-site.xml and move jar dependencies to HDFS for hive

```
<property>
  <name>hive.execution.engine</name>
  <value>spark</value>
  <description>Use Map Reduce as default execution engine</description>
</property>
<property>
  <name>spark.master</name>
  <value>spark://localhost:7077</value>
</property>
<property>
  <name>spark.eventLog.enabled</name>
  <value>true</value>
</property>
<property>
  <name>spark.eventLog.dir</name>
  <value>/tmp</value>
</property>
<property>
  <name>spark.serializer</name>
  <value>org.apache.spark.serializer.KryoSerializer</value>
</property>
<property>
  <name>spark.yarn.jars</name>
  <value>hdfs://localhost:54310/spark-jars/*</value>
</property>
```




Spark with Ceph

1. Add dependencies in pom.xml
 - a. aws-java-sdk 1.7.4
 - b. hadoop-aws 2.7.3
2. Configure spark-defaults.conf
3. Configure spark-env.sh
4. Setup hadoop configuration in Spark application
 - a. access key
 - b. secrete key
 - c. endpoint
 - d. NativeS3FileSystem



Spark with Ceph

Two steps:

- a. Creating ceph bucket and upload input files into ceph
- b. Using s3a to access data in Ceph for running spark jobs



Spark with Ceph

Creating ceph bucket and upload input files into ceph

```
client = boto3.client(service_name='s3', aws_access_key_id=access_key, aws_secret_access_key=secret_key,
                      endpoint_url=endpoint_url,
                      use_ssl=False,
                      verify=False,
                      config=Config(signature_version='s3v4'))

bucket_name = sys.argv[1]
direct_name = sys.argv[2]

file_name_listdir(direct_name)

client.create_bucket(Bucket=bucket_name)

for file in file_list:
    print(file)
    f = open(direct_name + file, 'r')
    data = f.read()
    client.put_object(Bucket=bucket_name, Key=file, Body=data)
```



Spark with Ceph

Using s3a to access data in Ceph for running spark jobs

```
sc = SparkContext()
sc._jsc.hadoopConfiguration().set("fs.s3.impl", "org.apache.hadoop.fs.s3native.NativeS3FileSystem")
#sc._jsc.hadoopConfiguration().set("fs.s3a.awsAccessKeyId", "0555b35654ad1656d804")
#sc._jsc.hadoopConfiguration().set("fs.s3a.awsSecretAccessKey", "h7GhXuBLTrlhVUyxSPUKUV8r/2EI4ngqJxD7iBdBYLhw
==")
#sc._jsc.hadoopConfiguration().set("fs.s3a.connection.ssl.enabled", "false")
sc._jsc.hadoopConfiguration().set("fs.s3a.awsAccessKeyId", "GC4A6H005IMZSTU3MNSA")
sc._jsc.hadoopConfiguration().set("fs.s3a.awsSecretAccessKey", "PkVehDJyYQMavZYWgrBFij9yieLnva1m9mOnelqx")
sc._jsc.hadoopConfiguration().set("fs.s3a.endpoint", "http://10.0.0.13:8000")
```

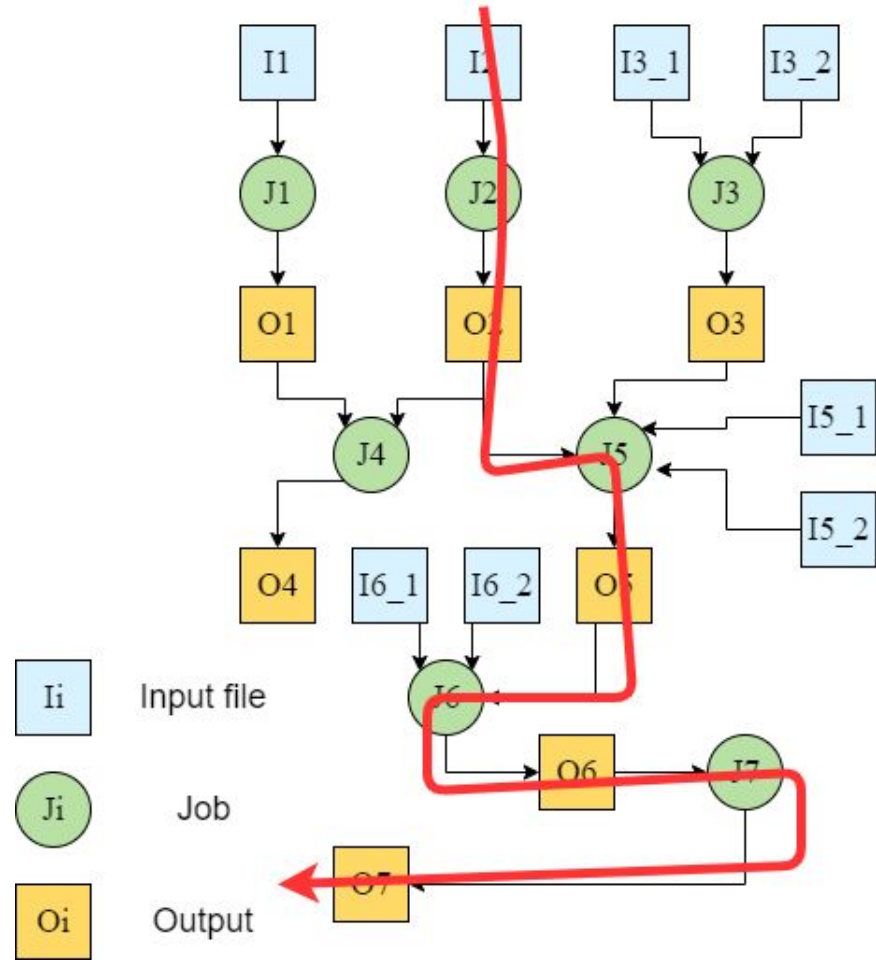


Spark with Kariz

1. Change the endpoint to Ceph RGW's endpoint
2. Kariz daemon uses port 3188 and cache daemon uses port 3187
3. Run these two servers
 - a. `${KARIZ_ROOT}/plans/kariz/api/server.py`
 - b. `${KARIZ_ROOT}/cache/server.py`
4. Run Spark application
 - a. Make sure Ceph is running correctly

Kariz Cache planner


- Find the path with longest running time by Dijkstra.
- Assuming the red line is the longest path, plan to prefetch I2, I5_1, I5_2, I6_1, I6_2
- Get data size and file location



Prefetch

Now we have all we need: what /when to cache, files on longest path, data size, kariz with spark, and kariz with ceph. Also : **Kariz already have the pig prefetching module.**

FIG: Next stage have to wait until all the node in last stage finish




Spark: No such stage dependency

Trick: fit the longest path into a new DAG


then use pig

Prefetching module.





Demo



Prefetched Data



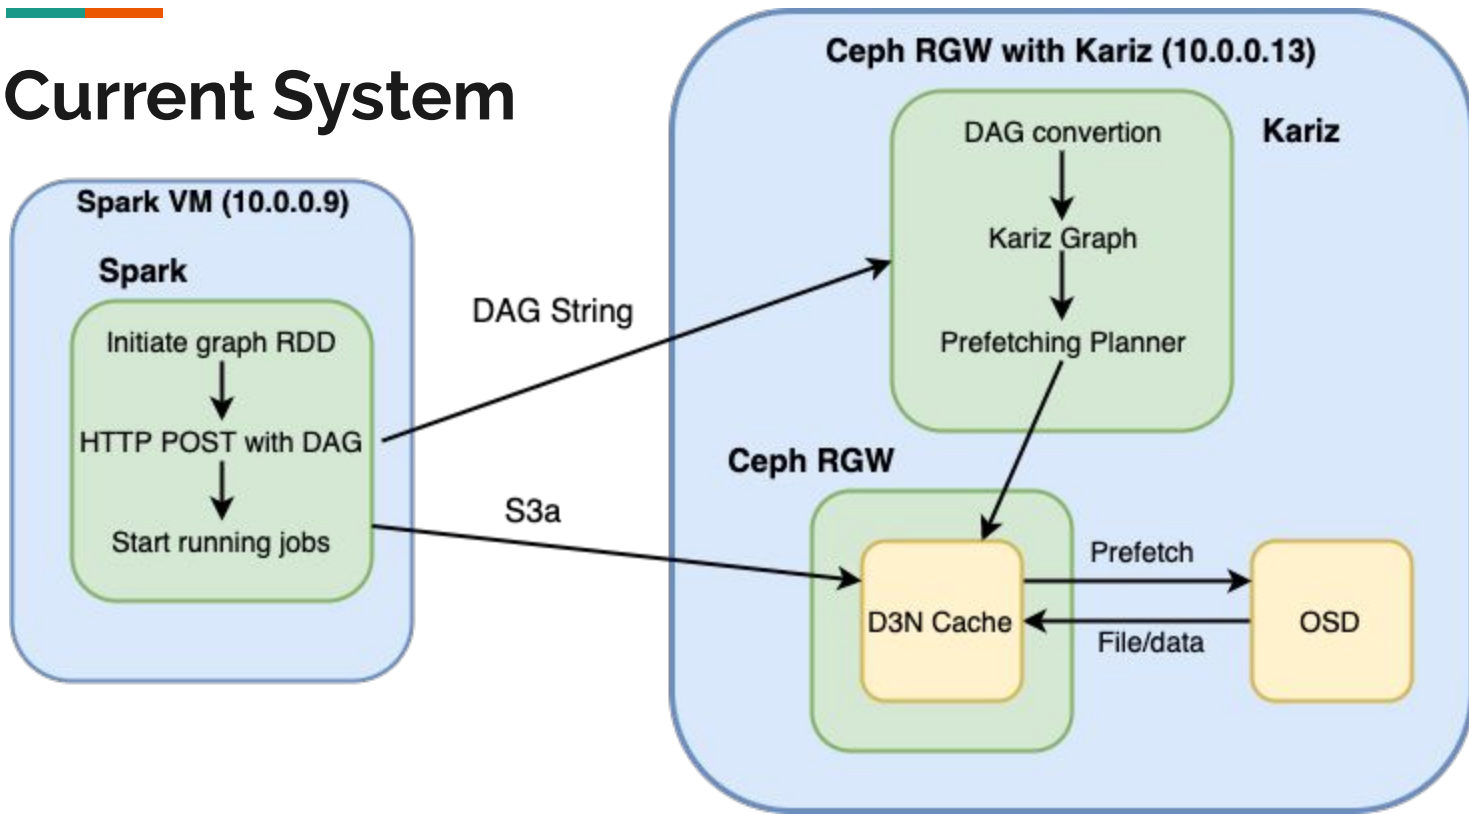
```
gangwei — centos@ceph-machine:/tmp — ssh 128.31.27.218 -i cloud-ceph-rg...  
[centos@ceph-machine ceph]$ cd /tmp  
[centos@ceph-machine tmp]$ ls  
3fe61bdb-9817-4bcc-97e5-49e0a5b56c44.4139.1__shadow_.7x1KsnKb03Ye10XHC8G-zMitozr  
Fhut_1  
3fe61bdb-9817-4bcc-97e5-49e0a5b56c44.4139.1__shadow_.jzWik70BS9Nvr0wgZNKA0628vNy  
LaBd_1  
ceph-asok.pPYP1G  
ceph-asok.WICVyy  
hadoop-centos  
systemd-private-6300acbf6f834ceda44f89613c21e30d-chronyd.service-PY3EQm  
[centos@ceph-machine tmp]$
```



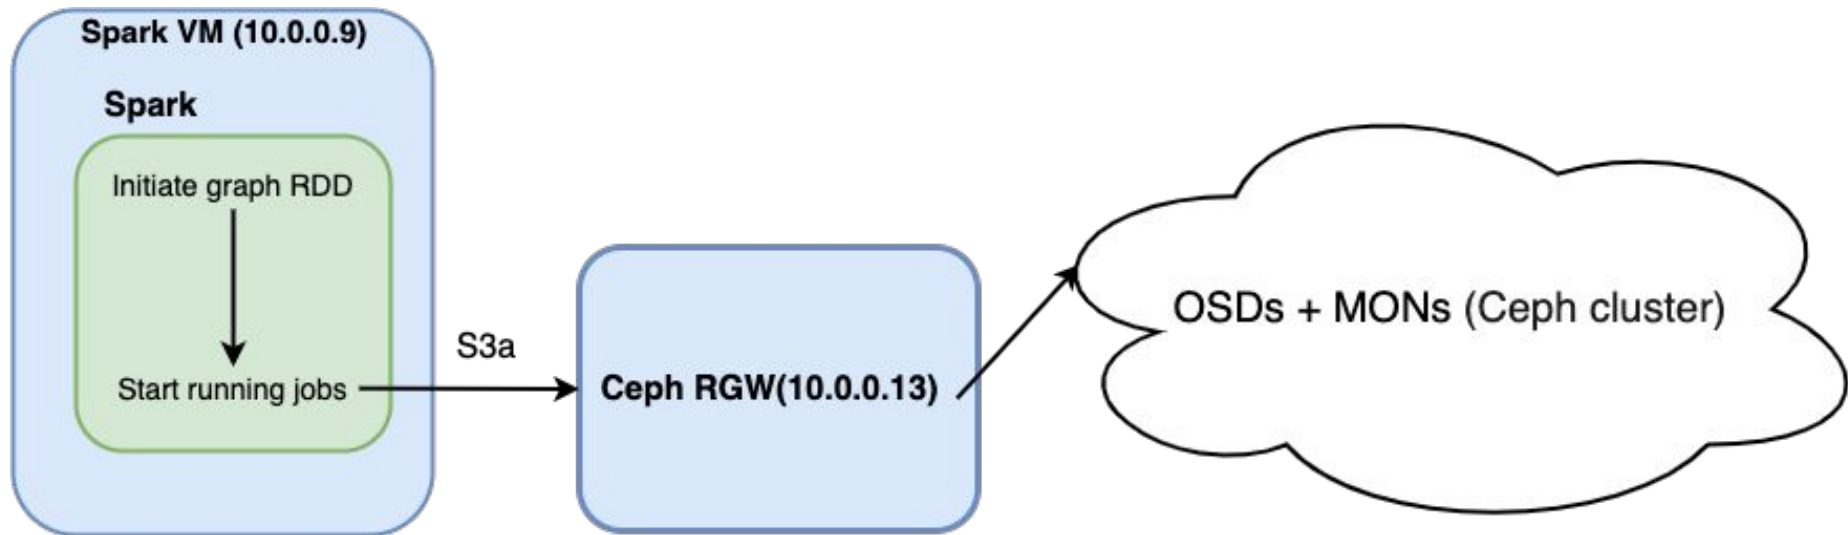
Next Step

1. Hive on Spark?
2. Performance Evaluation
 - a. System adjustment
 - b. Runtime comparison

Current System



Next Step



Next Step

